

DSCI 519: Foundations and Policy for Information Security

Theoretical Limits on System Security

Tatyana Ryutov

Reminders

- Complete Quiz 2 to by September 25, midnight
 - 45 minutes
 - Covers lectures 3 and 4
- Lab2 is posted, due by October 23, midnight
- Project proposal is due by October 2nd, midnight
 - If you decide to do the assigned project, you need to submit nothing

Outline

- Review
- Bell-LaPadula Multics interpretation
- General undecidability of security
 - Preliminaries
 - Halting problem
 - Turing machines
 - Safety question
 - Harrison-Ruzzo-Ullman result

L5.Q6



-
- Think about all the material covered in **lectures 1-5**
 - Indicate **specific** topics/questions for midterm review next lecture
 - focused questions/topics I can give you an answer to
 - If everything is clear, it's OK to say: “none”

Presentation 3

Multiple Independent Levels of Security Architecture (MILS)

Ashley Ma & Miliano Mikol

Presentation 4



Secure Hardware Extensions

Your Muddy Points

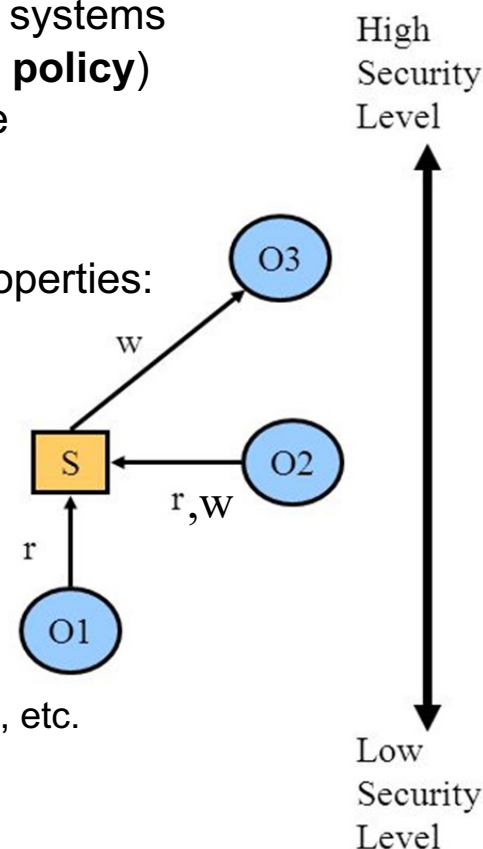
- What roles traditionally or formally are responsible for security mathematical formulas? Are the formulas derived during design/threat-modeling and then verified dynamically during integration and deployment?
- BLP's communication down solution - if a Trojan horse is embedded into the system while the subject has maximum security level, then it can copy data and when the security level downgrades, the Trojan can create a new document with lower security level with the previously copied data which had higher security clearance
- Not sure about how deep about basic security theorem we should understand. Just the concept or we should also understand the proof?
- Lattices, LUB and GLB
- BLP abstraction for system state

Outline

- Review
- **Bell-LaPadula Multics Interpretation**
- General undecidability of security
 - Preliminaries
 - Halting problem
 - Turing machines
 - Safety question
 - Harrison-Ruzzo-Ullman result

Review: BLP

- BLP a useful particularization of a class of systems
- The goal is to enforce confidentiality policy
- Methodology
 - Define an abstract **model** that can be used to describe computer systems
 - Define what it means for a system in the model to be secure (**the policy**)
 - Develop techniques to prove that a system in the model is secure
- Approach:
 - Use state-transition systems to describe a computer system
 - Define system as secure IFF every reachable state satisfies 3 properties:
 1. simple-security property: S can read O IFF $S \text{ dom } O$
 2. *-property: S can write O IFF $O \text{ dom } S$
 3. discretionary-security property
 - Prove a Basic Security Theorem (BST)
- Main contributions:
 - The overall methodology to show that a system is secure
 - adopted in many later works
 - The state-transition model
 - which includes an access matrix, subject security levels, object levels, etc.
 - The introduction of *-property
 - ss-property is not enough to stop illegal information flow



Review: BLP Abstraction for System State

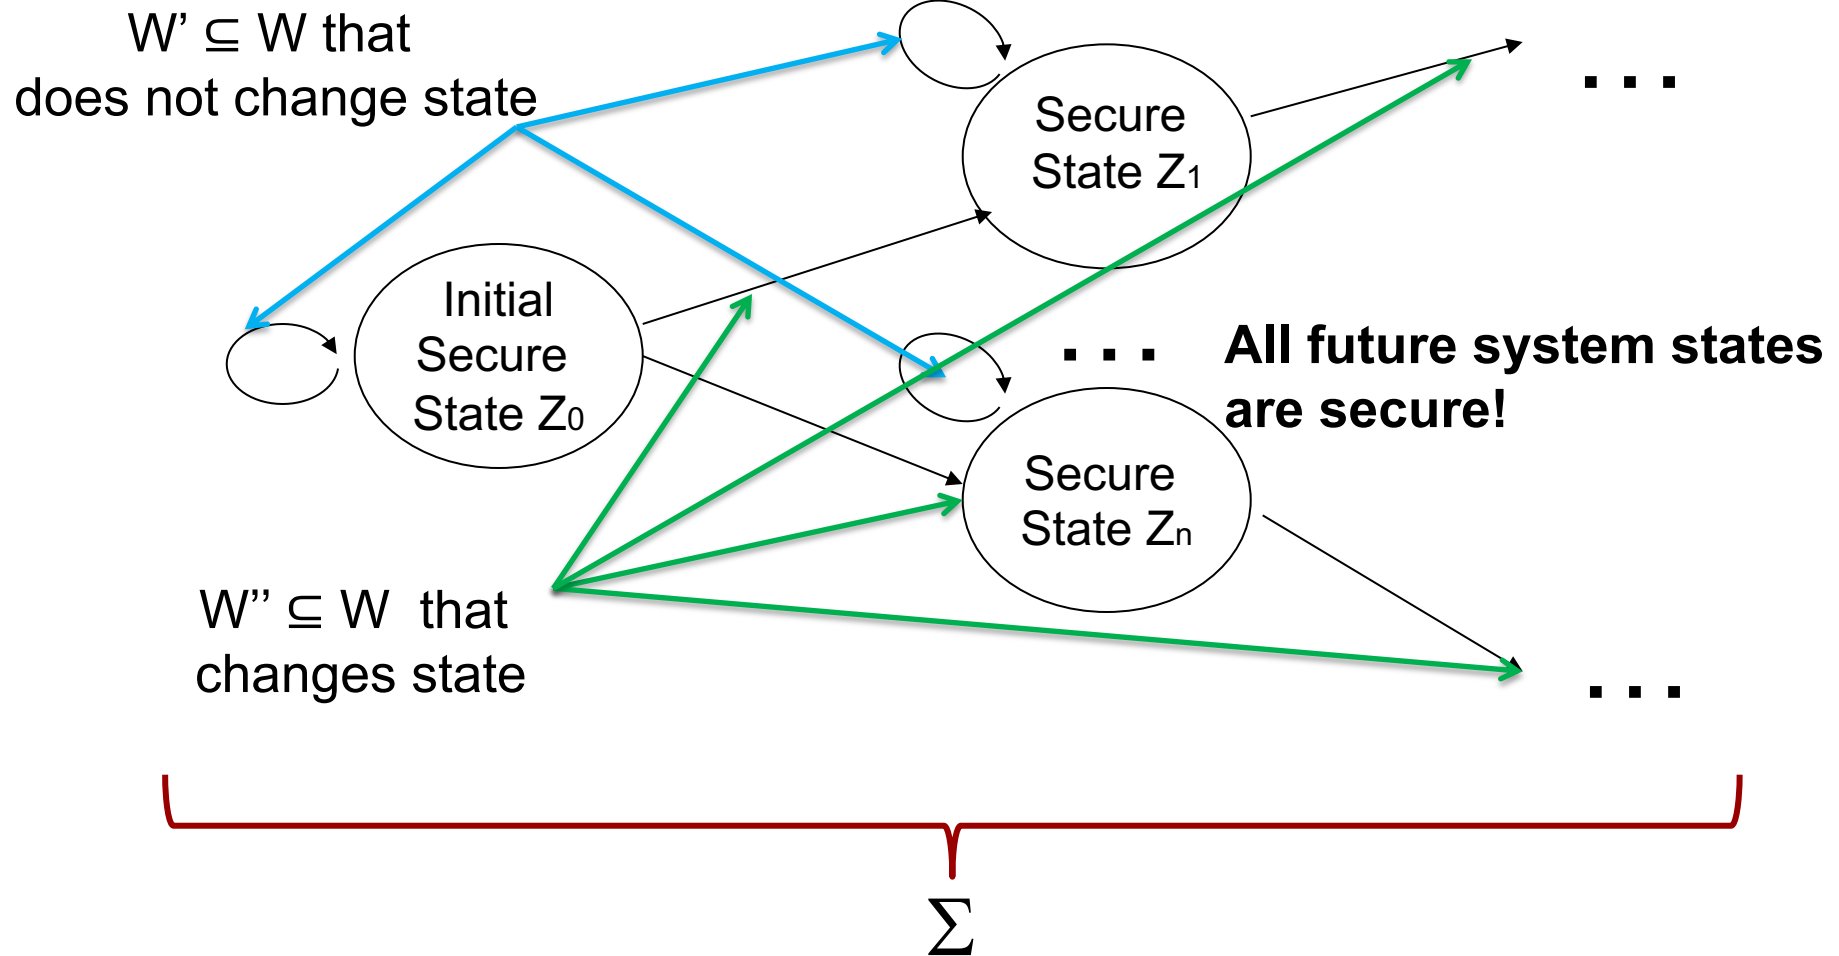
- The **state** of the system is (b, M, f, H) where:
 - **b** indicates which subjects can access which objects
 - The set of rights that may actually be exercised in the **current** state
 - **M** is the access control matrix for the current state
 - DAC
 - MAC can make some rights unusable
 - **f** is tuple indicating subject and object access classes
 - MAC
 - For subjects: max and current clearances
 - **H** is the *hierarchy* of objects (for naming objects)
 - Can not do access control unless can uniquely name objects
- Recall that we have defined a “secure” state
 - MAC (Simple security, *-property) and DAC

b is a subset of all accesses that a subject can have in the system based on DAC (*M*) and MAC, why here it is subset but not the exactly the same set?

Review: System Representation

- System is represented by $\Sigma(R, D, W, z_0)$ where:
 - R denotes the set of requests for access (inputs)
 - D denotes the set of outcomes (outputs)
 - (y)es, (n)o, (i)llegal, (o) error
 - W is the set of actions of the system
 - At any time t the system is in state $v \in V$, where V is set of all system states
 - Recall that the $v = (b, m, f, h)$
 - W moves system from a state in V to another (possibly different) state in V
 - Each action is of the form $((r, d, (b, m, f, h), (b', m', f', h'))$
 - Example of W – kernel calls
 - z_0 is the initial state of the system
- System is all sequences of <request, decision, state> triples, with initial state z_0

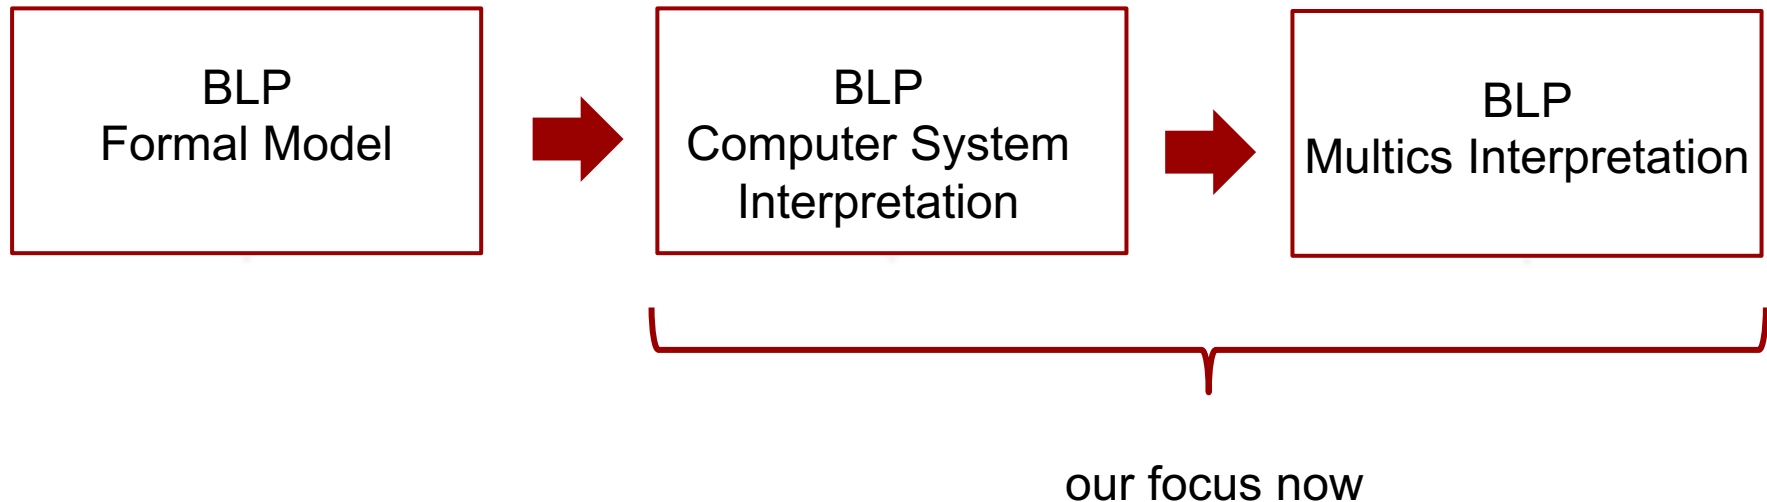
Review: BLP as a State Machine System View



- Define **all possible** transitions (W), each transition satisfies the 3 properties

W is described as the set of actions of the system. Does it mean transitions between states?

BLP Interpretation



MULTICS System

- Multiplexed Information and Computing Service (MULTICS)
 - Joint project by MIT and Bell Telephone Laboratories
 - Extraordinarily influential early time-sharing operating system
 - Unix, Linux, Paging, Segmentation, Protection, etc. had their birth in this project
- MULTICS was commercialized 1973-1985, in use until 2000
 - Honeywell commercial product beginning ~ 1970
- Designed for security from outset
 - But repeatedly broken into: penetrate & patch
 - Current industry “best practice”
 - Eventually led to redesign based on FSPM and RM
- Proof of performance/usability of secure systems



Multics Correspondence to BLP State

b → segment descriptor words (SDW)

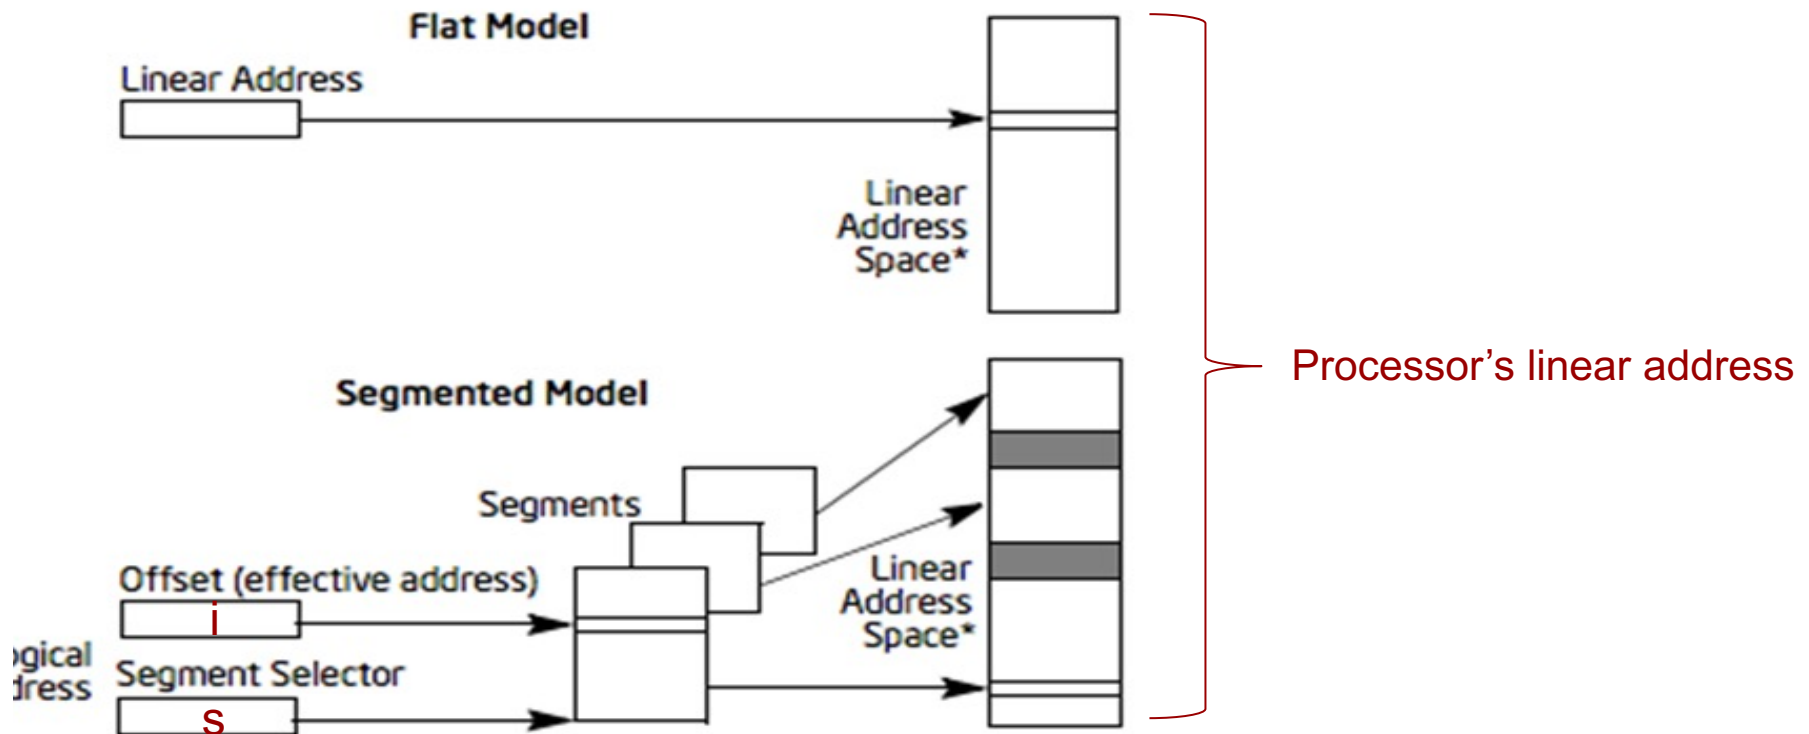
M → access control lists (ACL)

f → $\left\{ \begin{array}{l} \text{segment security level in directory} \\ \text{process security level tables} \end{array} \right.$

H → branches in a directory

Sidebar: Flat vs. Segmented Memory

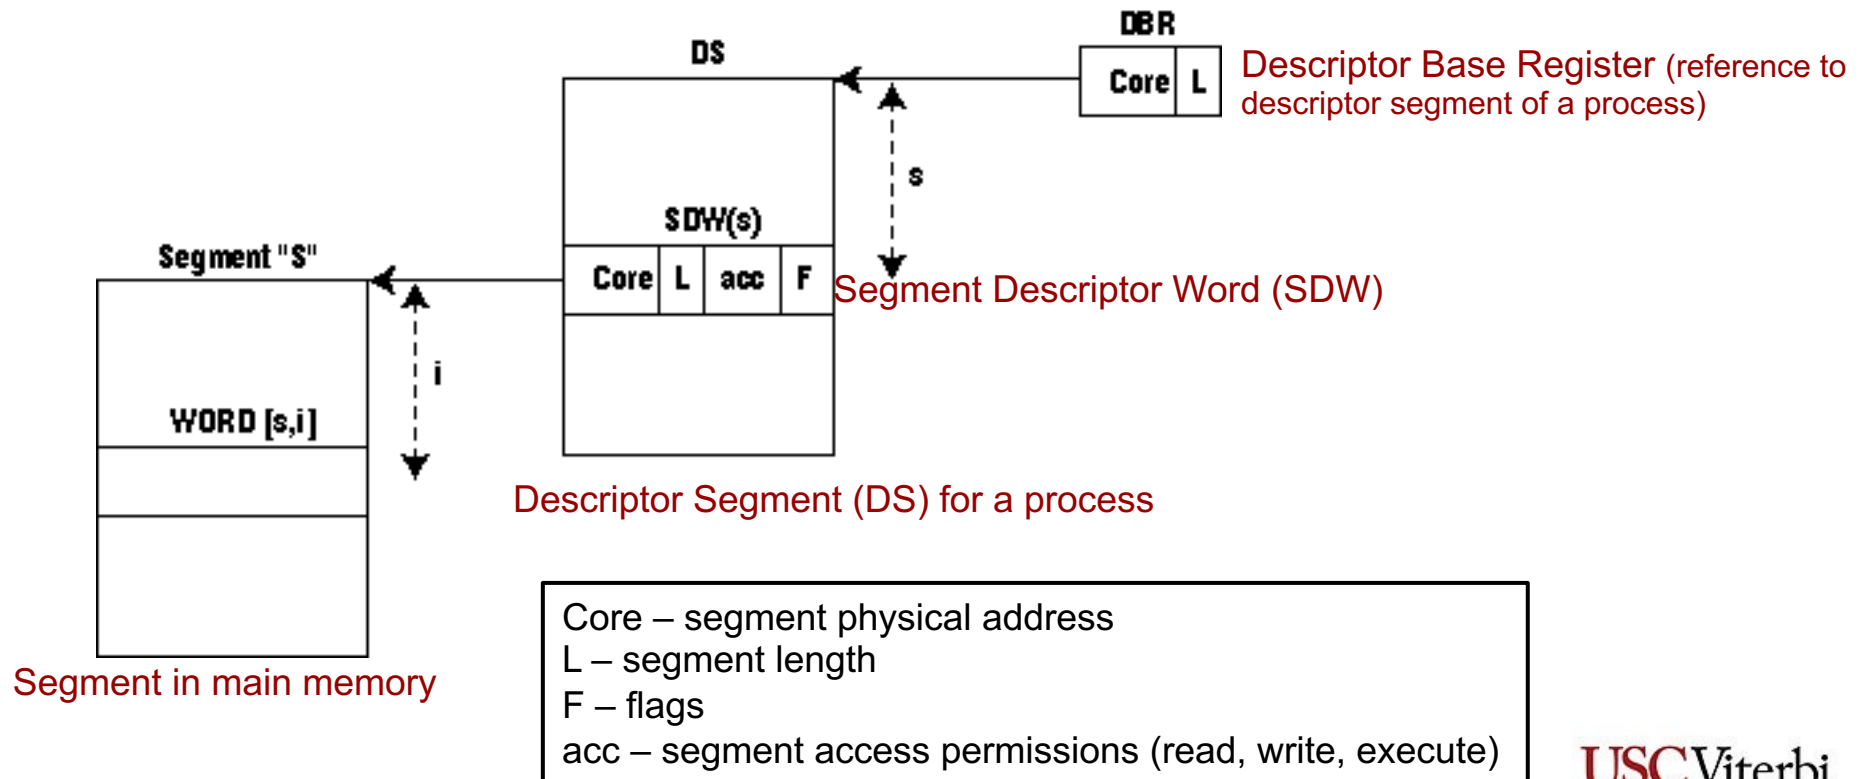
- Segmented memory model: memory appears to a program as a group of independent address spaces called **segments**



Virtual Address of a segment = $[s, i]$

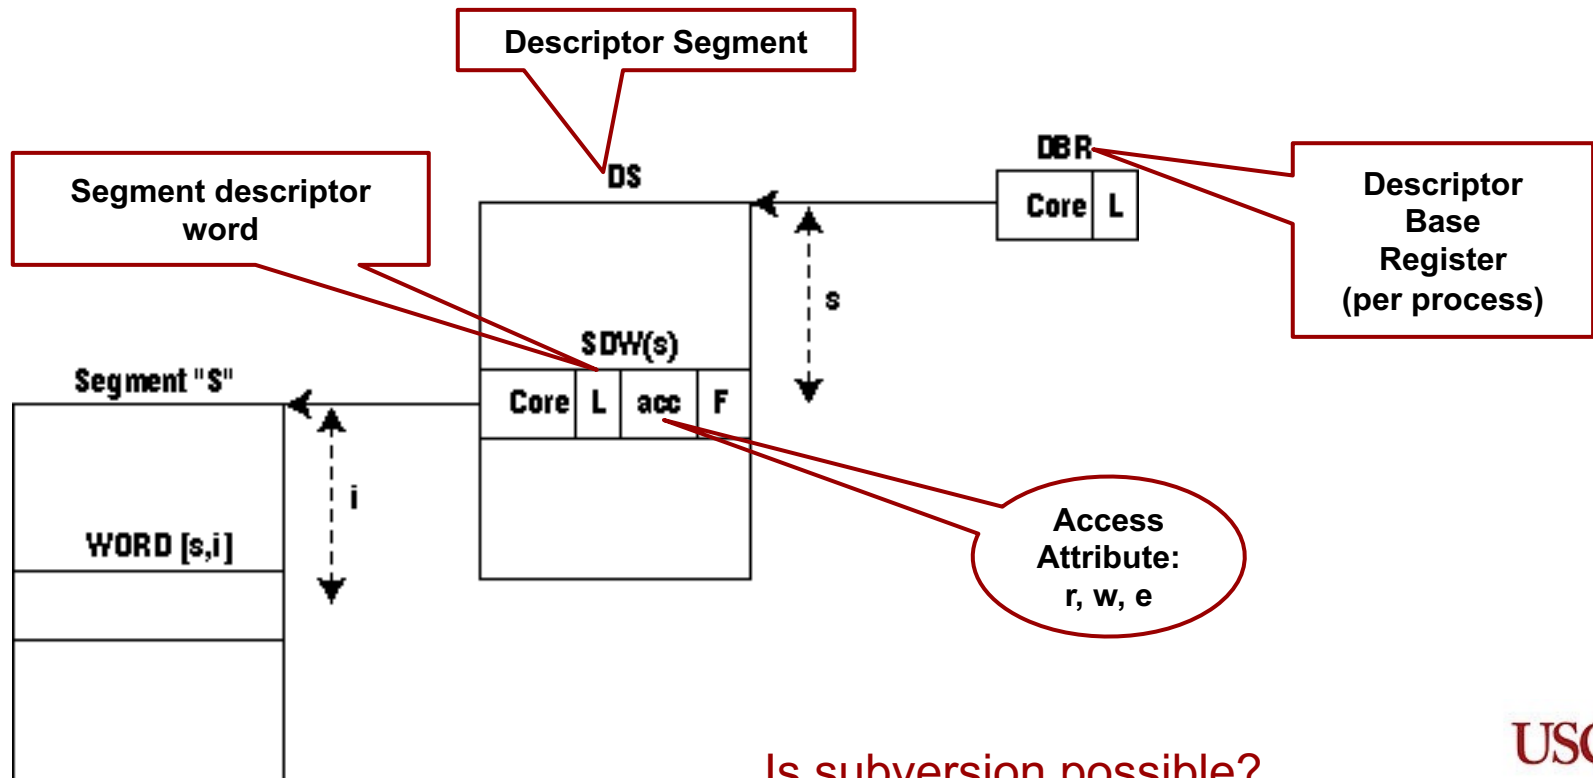
Multics Hardware Memory Referencing

- All code, data, I/O devices, etc. that may be accessed by a process are stored as segments
- A segment virtual address consists of a pair of integers [s, i]
 - "s" is called the segment number (descriptor)
 - "i" the index within the segment (offset)
- Descriptor Segment (DS) stores Segment Descriptor Words (SDW) that reference each of the process's active segments



State for Hardware Current Access (b)

- The set **b** is interpreted as a descriptor segment (DS)
- The set **b** has elements (S, O, p) – interpreted as a SDW
 - Subject S has current access to object O in mode p
- Each subject S is a distinct process
- Each distinct data segment is an object O



Is subversion possible?

Multics Correspondence to BLP State


b → segment descriptor words (SDW)

M → access control lists (ACL)

f → $\left\{ \begin{array}{l} \text{segment security level in directory} \\ \text{process security level tables} \end{array} \right.$

H → branches in a directory

Multics Hierarchy Interpretation




Relationship H between objects:
Both data and directory segments stored in a “branch” of directory

Multics Correspondence to BLP State

b → segment descriptor words (SDW)

M → access control lists (ACL)

f →  segment security level in directory
process security level tables

H → branches in a directory

ACL: Interpret Access Matrix

Objects Subject	Segment 1	Segment2	Segment 3	Segment4
Bob Process	read	read, write		write
Flo Process	read, write	write		
Alice Process	read	read	read	read, write
Dan Process	read		read, write	read

ACL for Segment 1:

Bob	read
Flo	read, write
Alice	read
Dan	read

Access Control List (ACL):

- A column of Matrix (M)
- Stored in a “branch” of directory
- Branch contains physical location and ACL for the object

Multics Correspondence to BLP State

b → segment descriptor words (SDW)

M → access control lists (ACL)


f → {
 segment security level in directory
 process security level tables

H → branches in a directory


Interpret Security Level Function (f)

- Segment Access Class and ACL are stored in parent directory of the segment
- Process Access Class is stored in TCB (internal kernel table)

Directory Segment



Security Label for
a segment



TCB internal table

Bob	Top Secret
Flo	Top Secret
Alice	Unclassified
Dan	Unclassified

Security Label
for a process

Interpretations of RM Access

- Use access modes to describe access control
 - Must interpret access mode
- Fundamentally, access devolves to read & write
 - Memory chips have only read and write functionality
- Other derivative modes ultimately map to that

Model Multics access rights

r → read

r,e → execute




w → read and write

a → write

Recall: BLP System Representation

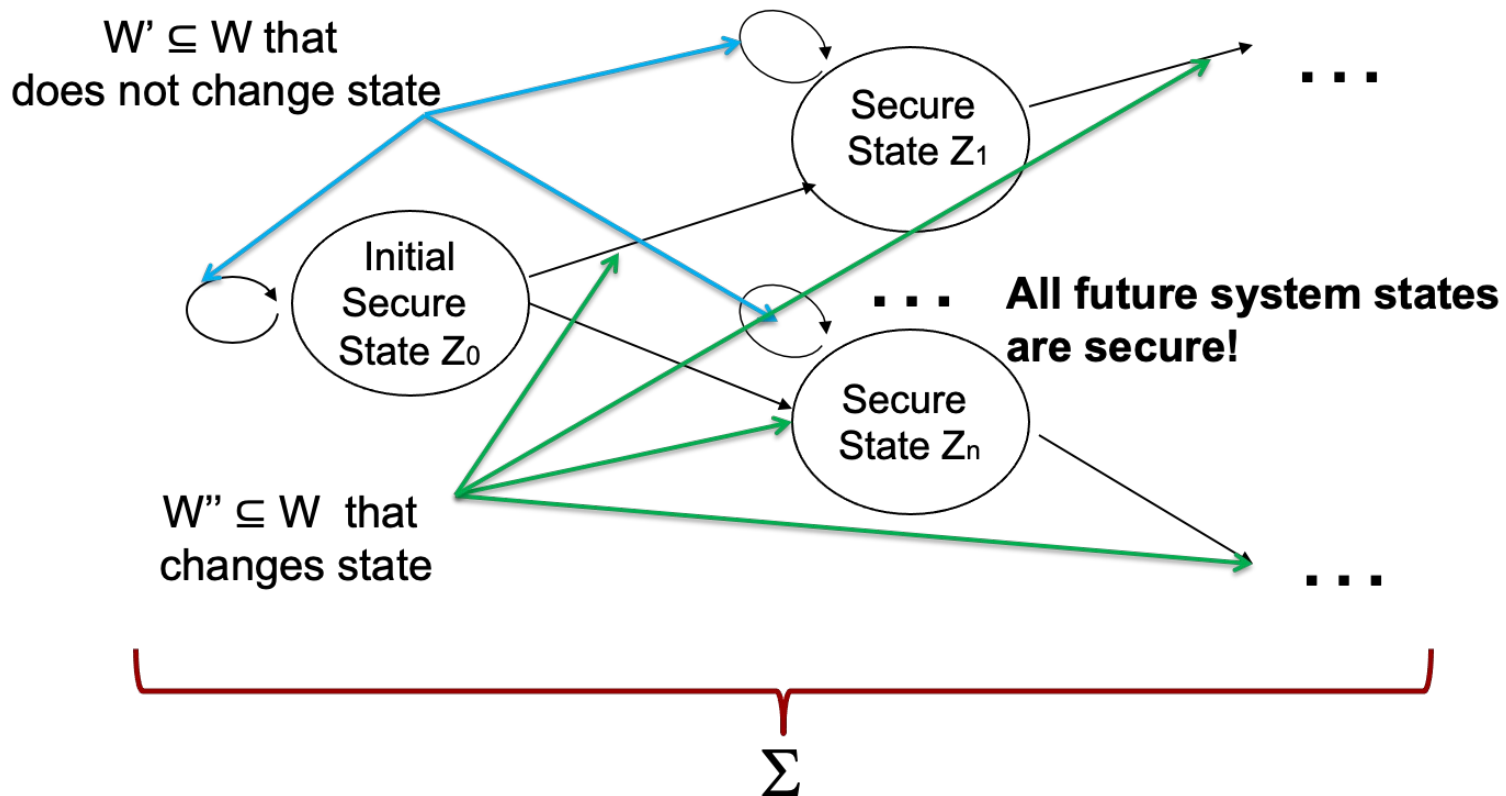
- Represented by an initial state and a **sequence** of request, decisions, and states
- In formal terms, system $\Sigma(R, D, W, z_0)$ where:
 - R denotes the set of requests for access
 - D denotes the set of outcomes
 - W is the set of actions of the system
 - z_0 is the initial state of the system
- **W** notation in formal terms:
 - $W \subseteq R \times D \times V \times V$ is the set of actions of system
 - System moves from a state in V to another (possibly different) state in V
 - Usually “yes” means state change, but if segment is already in b , no change in state
 - Each action is of the form $((r, d, (b, m, f, h), (b', m', f', h'))$
 - Recall that the set (b, m, f, h) is a state of the system

Basic Security Theorem

- $\Sigma(R, D, W, z_0)$ is a secure system if
 - z_0 is a secure state
 - Every action $(r, d, (b, m, f, h), (b', m', f', h'))$ in W satisfies:
 - a. Every $(s, o, p) \in b - b'$ 
 - simple security condition 
 - *-property
 - ds-property
 - b. Every $(s, o, p) \in b'$ not satisfying the following, not in b 
 - simple security condition
 - *-property
 - ds-property
- i.e., any **new** accesses in b must satisfy the policy*
- i.e., must revoke **old** accesses that would no longer satisfy the policy (due to changes in subject level, or DAC permissions)*

What is missing?

- Given a system, can you show that **every** action (usually means kernel calls and their effects) maps to an FSPM transition (a member of set W)?



- Define **all possible** transitions (W), each transition satisfies the 3 properties

General Mechanism of Model

- “Inductive nature” of security
 - Each alteration of state does not breach security
 - If preserved from one state to next, system is secure
 - Total system security guaranteed systematically
 - For every access tuple (s,o,p) added to current access set b
 - Meets ss , $*$ -property, and ds -property
- Modular specification of system capabilities
 - Only **one** rule for each (request, current state) pair
- Goal: relate rule properties to system properties
 - Preservation of properties over **one transition**

Recall: Reference Monitor (RM) Functions

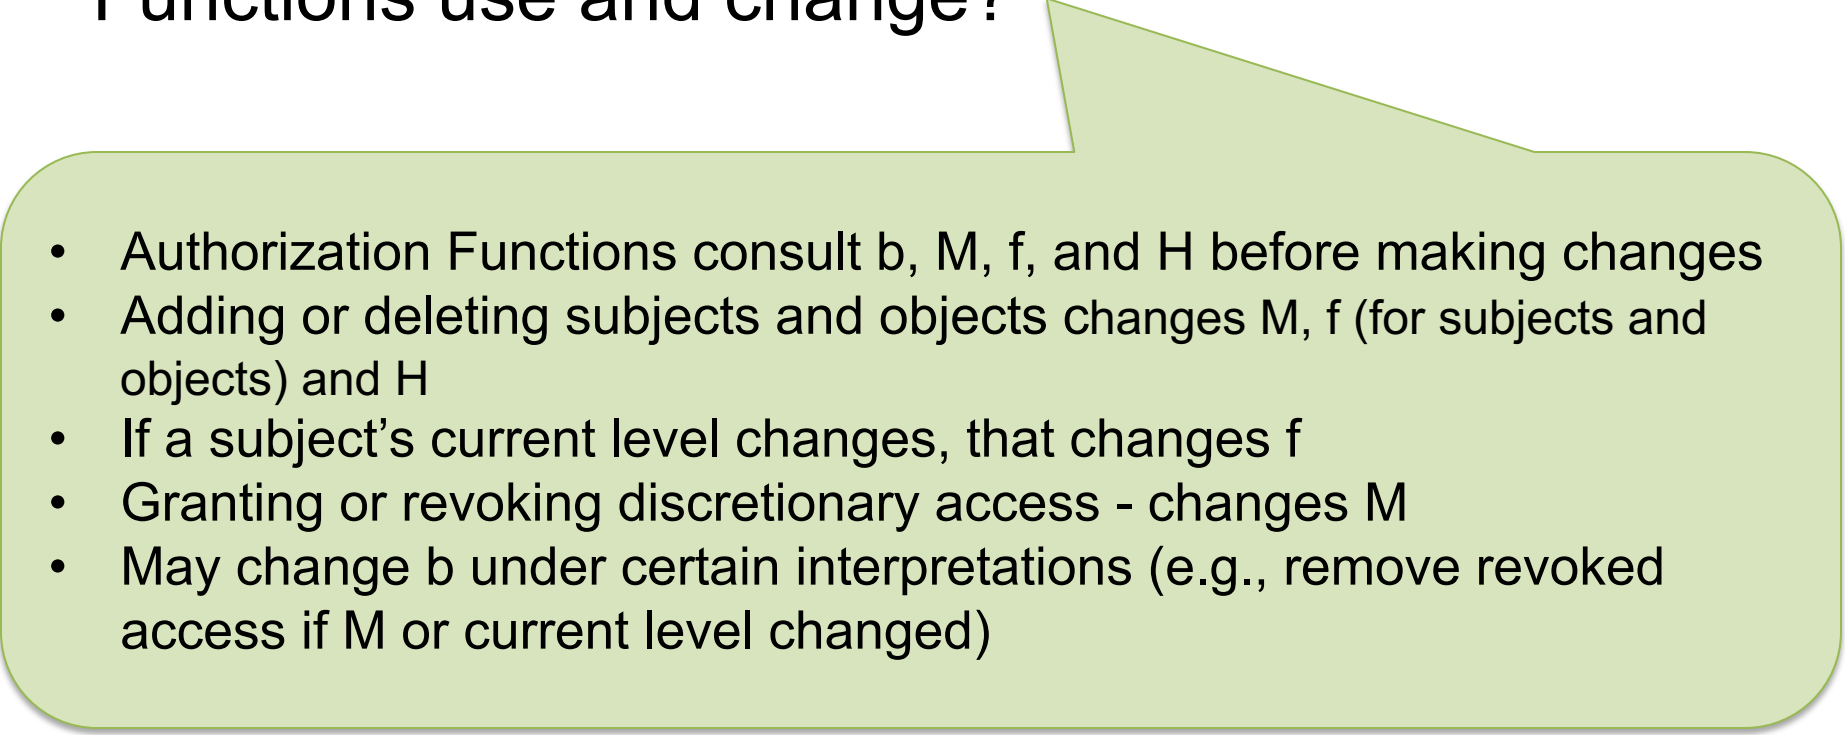
- Two classes
 - Reference and Authorization
- Reference functions
 - Subject wants to access an object – is it OK?
- Authorization functions
 - Change in authorization data base
 - E.g., add/remove user or access right

Reference Functions in BLP

- What state elements (b, M, f, H) do the RM Reference Functions use and change?

- Reference Functions consult b, M, and f before making changes
- Check b (enforcement in hardware), b acts as “cache” for granted access but in hardware
- If not authorized by b, check f to see if the MAC policy permits access
- If MAC permits, the last step before changing b is to consult M
- If access is granted, b will be changed

Authorization Functions in BLP

- What state elements (b, M, f, H) do Authorization Functions use and change?
- 
- Authorization Functions consult b, M, f, and H before making changes
 - Adding or deleting subjects and objects changes M, f (for subjects and objects) and H
 - If a subject's current level changes, that changes f
 - Granting or revoking discretionary access - changes M
 - May change b under certain interpretations (e.g., remove revoked access if M or current level changed)


Four Classes of General Functions

- Four classes of functions that can change system state:
 1. Functions to alter the current access (set b)
 2. Functions to alter current access permission (M)
 3. Functions to alter the level functions (set f)
 - Level (subject)
 - Level (object)
 - Current-level (subject)
 4. Functions to alter object structure (hierarchy H)
- **Next:** instantiate model rules in kernel primitives
 - Use notation in BLP Multics Interpretation

Morphisms from Multics to Model

- Instantiate BLP rules as 11 **specific** kernel primitives
 - Rule numbers are from BLP model
- 1. Alter current access (set b) \rightarrow SDW
 1. Kernel function: get-read
 2. Kernel function: get-write-only
 3. Kernel function: get-execute
 4. Kernel function: get-read-write
 5. Release read/execute/write
- 2. Alter access permission (matrix M) \rightarrow ACL
 6. Give read/execute/write
 7. Revoke read/execute/write

A **morphism** is a structure-preserving mapping from one mathematical structure to another




Morphisms from Multics to Model

3. Alter hierarchy (set H) → branch in directory

- 8. Create-object
- 9. Delete-object-group

4. Alter level function (set f)

- For subject → current-security-level table
- 10. Change-subject-current-security-level
- For object → branch in directory
- 11. Change-object-security-level




In practice, not used due to the tranquility principle

Paraphrased Multics BST

- Basic Security Theorem – Multics (Σ) is secure
 - If z_0 is a secure initial Multics state and
 - All Multics kernel primitives (rules for actions in set W)
 - Satisfy the simple security condition
 - Satisfy the *-property
 - Satisfy the discretionary security property
- ***Means all future Multics states are secure***

Multics as a State Machine

Multics kernel calls
do not change state



- Defined **all possible** transitions (Multics kernel calls), each transition satisfies the 3 properties

Other BLP Interpretation Experience

- Legacy of past security kernel deployment
 - Different deployments over periods of 30-40 years, no security patches
- **Class A1** (formally verified) Security Kernel worked examples:
 1. SACDIN - Strategic Air Command Digital Network
 - Minuteman missile control (IBM)
 - LGM-30 Minuteman is a U.S. land-based intercontinental ballistic missile
 2. SCOMP for Multics comms (Honeywell)
 - Secure front-end communications processor for Multics
 3. Secure Ethernet LAN (Boeing)
 4. GTNP/GEMSOS (Gemini Computers, Inc.)
 5. Blacker “VPN” front-end (Unisys for NSA)
 - End-to-end encryption on the United States' Defense Data Network
 6. VAX secure VMM - not shipped (DEC)
 - Developed with a heavy emphasis on performance and on system management tools



Properties of FSPM of RM Abstraction

- Only known high assurance cyber security basis
 - Provides the basic theory of computer security
 - Essential for trustworthy response to subversion
- Also useful to assess low to medium assurance
 - When fails to meet all three RM principles
 - E.g., improving an OS
 - Improve” OS, not necessarily make it “secure”
 - E.g., RM functions in an application program
 - E.g., a DBMS
 - But note that it is still on low-assurance OS

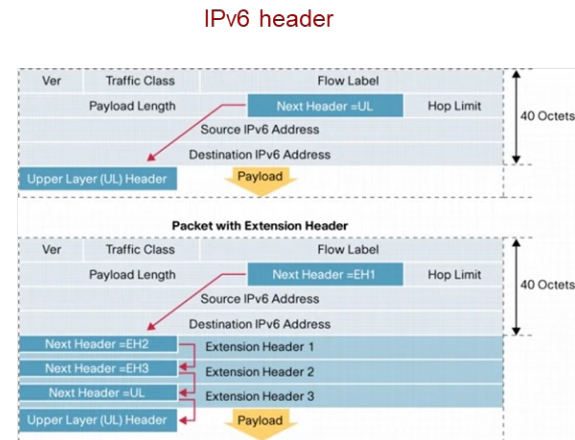
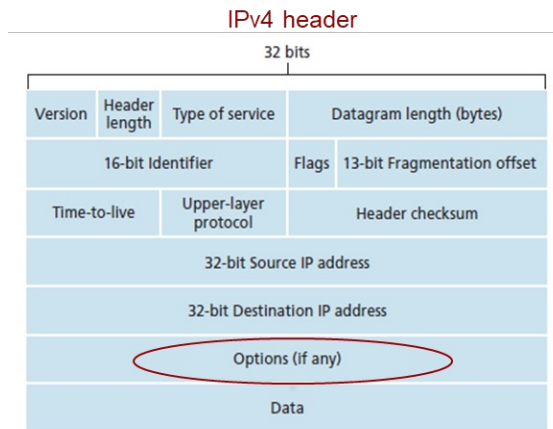
Case Study: Target Breach



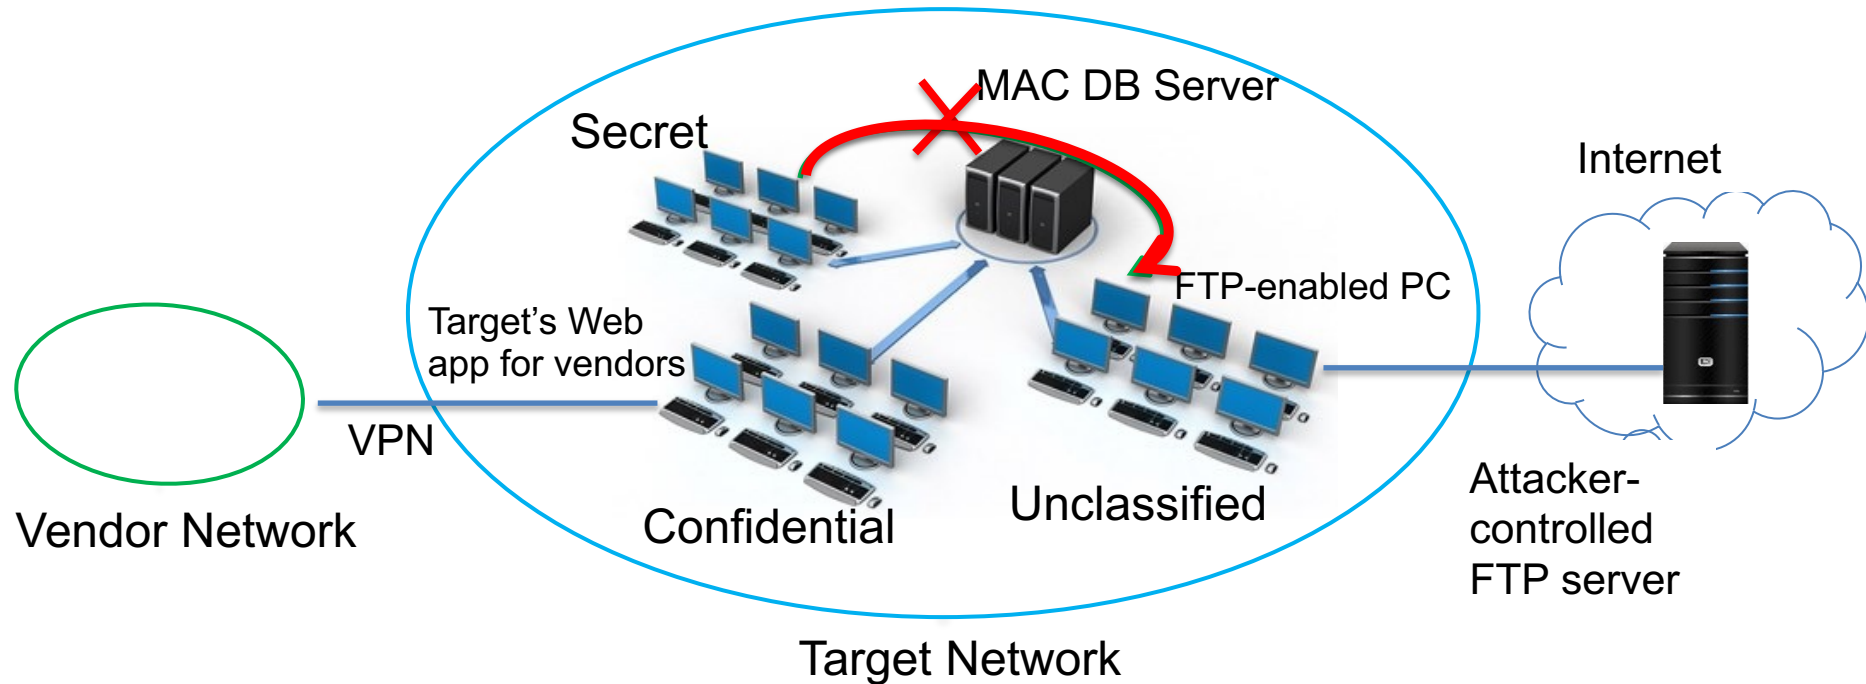
- 2013 Target breach: 70M customers' personally identifiable information (PII) and 40M credit cards stolen (CC)
 1. Initial penetration through stolen vendor's credentials
 2. Gain access to a Target web services for vendors
 3. Deploy malware on many Target's POS machines which were used to steal credit card info
 4. Sent stolen credit cards periodically to a central repository within Target's network using standard Windows protocols
 5. Exfiltrated stolen data from the central repository to the attackers' controlled server via FTP

Proposed Solution for the Target Case

- Use BLP to enforce information flow
- Assign access classes to data:
 - PII data – (Secret, {Account})
 - CC data – (Secret, {Payment})
 - Vendor data – (Confidential, {HVAC})
 - Public data: Unclassified
- MLS DBs for PII and CC
- NTCB concept – partition network into TCB subsets (covered later in the course)



MAC Solution for the Target Case



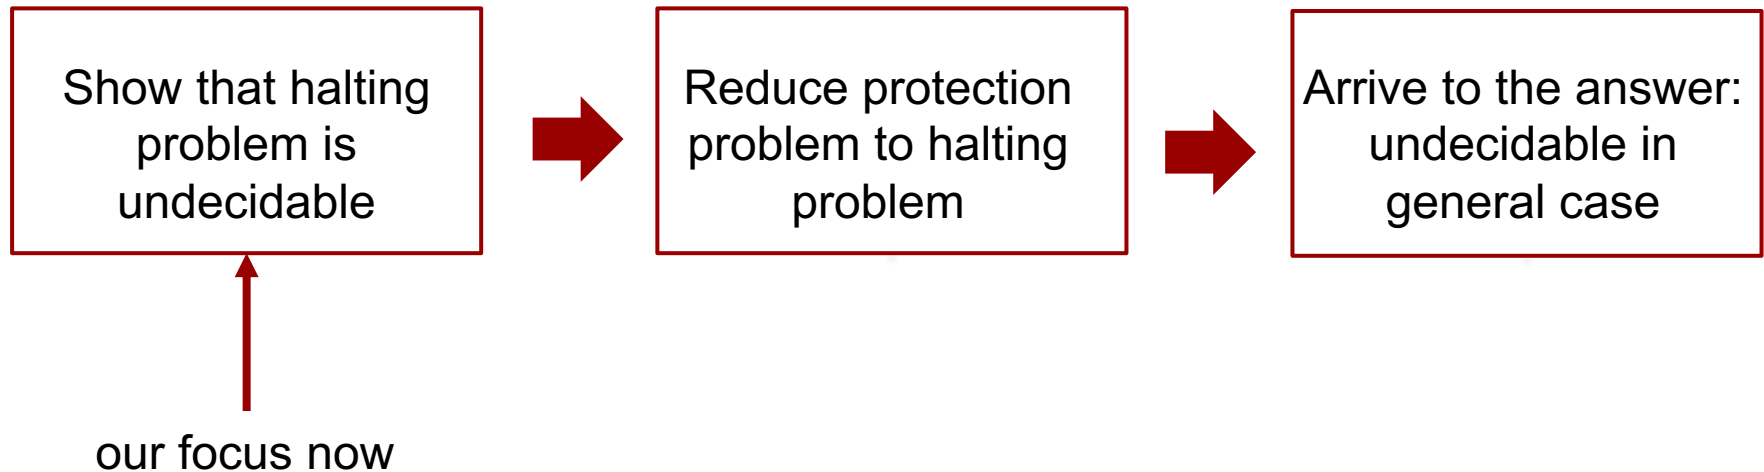
Can't exfiltrate sensitive data!

Outline

- Review
- Bell-LaPadula Multics Interpretation
- General undecidability of security
 - Preliminaries
 - Halting problem
 - Turing machines
 - Safety question
 - Harrison-Ruzzo-Ullman result

What are we trying to do?

- Can we determine if a computer system is “secure” (i.e., implements our policy with high assurance)?
- Our goal is to answer the key question: is this protection problem decidable?

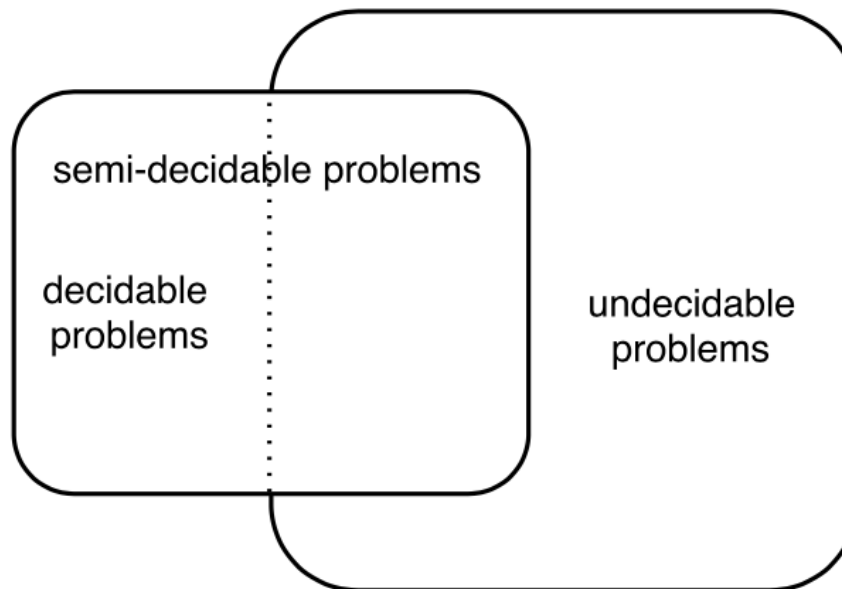


General Undecidability of Security

- Is there a generic algorithm to determine if computer system is “secure”?
 - This is the core question for information assurance
- Foundational results from 1976 paper
 - Harrison, Michael A., Walter L. Ruzzo, and Jeffrey D. Ullman. "Protection in operating systems." Communications of the ACM 19.8 (1976): 461-471
 - Based on Turing machines and undecidability
 - Based on reduction to “halting problem”

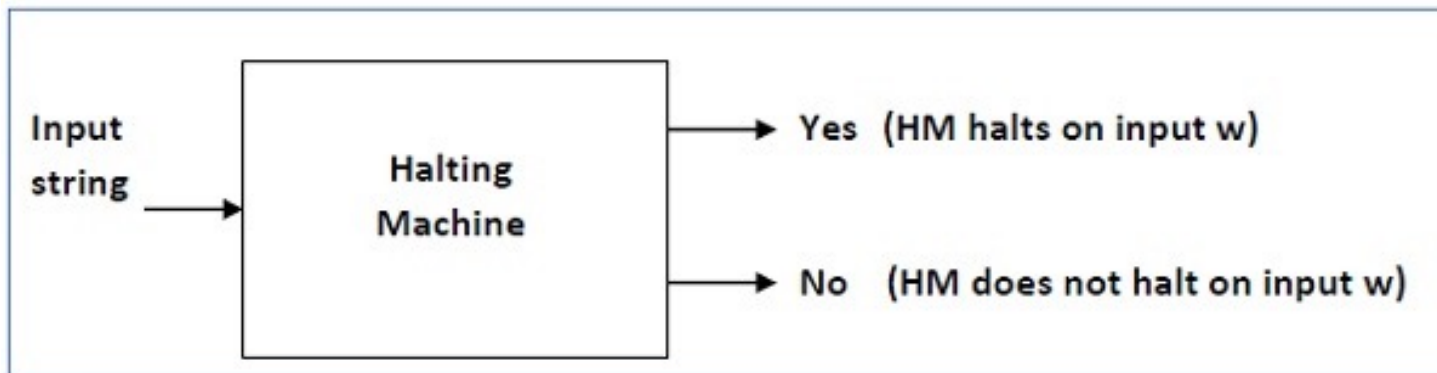
Decidability of Decision Problems

- Decidable problem
 - A decision problem can be solved by an algorithm that halts on **all** inputs in a **finite** number of steps
- Undecidable problem
 - A problem that cannot be solved for **all** cases by **any** algorithm
 - For a problem to be undecidable we just have to prove that there is **one** case it cannot produce an answer for



Halting Problem

- Is there a function (algorithm) that can look at **any** computer program and **any input** to that program and decide if the program will halt (not run infinitely)?
 - Question posed by Alan Turing in 1936 to prove that there are unsolvable problems



Proof by Contradiction


- Proof by contradiction: something that leads to a contradiction can not be true, and if so, the opposite must be true
 - To prove a statement by contradiction:
 - Start by assuming the **opposite** of what you want to prove
 - Then show that the consequences of this premise are impossible
 - This means that your original statement must be true
 - Example: prove that there is no largest number

Halting Problem: Proof

- Is there a function (algorithm) that can look at any computer program and any input to that program and decide if the program will halt?
- Assume existence of such function:
 HALT (program, input) → Boolean, where
 - **program** is a code of a program we want to test
 - **input** is the input to the program we want to test
 - **output**: **true** if the program halts with this input, **false** otherwise
- Now define a new function:
 SELF-HALT (program)
 {
 if (HALT(program, program))
 infinite loop
 else
 halt
 }

Halting Problem: Proof

- What if we use **SELF-HALT** to analyze itself?
 1. **SELF-HALT (SELF-HALT)** loops forever if **HALT (SELF-HALT, SELF-HALT)** is true
 - So if **SELF-HALT** halts on itself, it loops forever. **A contradiction!**
 2. **SELF-HALT (SELF-HALT)** halts if **HALT (SELF-HALT, SELF-HALT)** is false
 - So, if **SELF-HALT** does not halt on itself, it halts. **A contradiction!**
- This contradiction is unavoidable thus proving that the halting problem is undecidable




```
SELF-HALT (program)
{
    if (HALT(program, program) )
        infinite loop
    else
        halt
}
```

Importance of The Halting Problem

- This is the first computer problem proven to be **undecidable** (unsolvable)
 - Impossible to construct a **single** algorithm that always leads to a correct yes/no answer to a decision problem on infinite set of inputs
- To show that a problem is undecidable, it is sufficient to show that it is **equivalent** to the halting problem
 - More precisely, it is sufficient to show that if a solution to the new problem existed, it could be used to solve the halting problem, which is known to be undecidable

Turing Machine (TM)

- A Turing machine is a **theoretical** computing machine invented by Alan Turing in 1936 to serve as an idealized model of computing
- TM is a finite state machine with an infinite tape from which it reads its input and on which it records its computations one cell at a time
- TM can perform a special action – it can stop (halt)
- Despite its simplicity, the machine can simulate ANY computer algorithm, no matter how complicated it is!




TM: Informally

- TM takes input symbol and according to it and the current state does the following:
 1. it reads the symbol on the tape and writes another (new or the same) symbol on the tape
 2. moves the tape right or left by one cell
 3. changes to a new state

TM: Simple Example

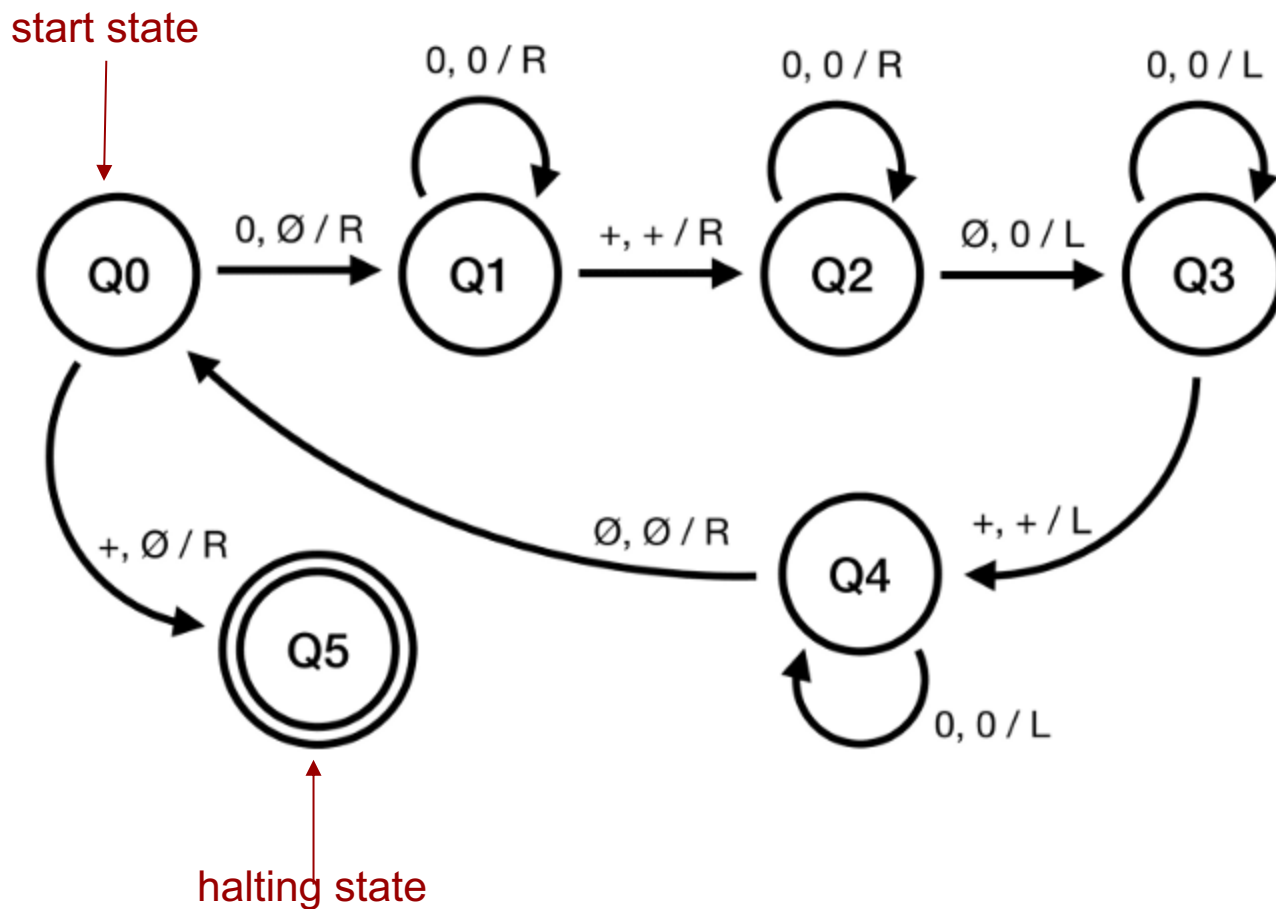
- Task: convert the 1s to 0s and vice versa

<i>State</i>	<i>Symbol read</i>	<i>Write instruction</i>	<i>Move instruction</i>	<i>Next state</i>
<i>State 0</i>	Blank	None	None	<i>Stop state</i>
	0	Write 1	Move the tape to the right	<i>State 0</i>
	1	Write 0	Move the tape to the right	<i>State 0</i>




TM: Another Example

- Task: perform addition, given input “00+000” (2 + 3) output “000000” (5)
- Use four symbols: “0”, “1”, “+” and “ \emptyset ”
 - “ \emptyset ” represents a blank space



TM: Another Example

1. Replace the first 0 with a blank space
2. Move to the end of the string
3. Add a "0" to the end of the string
4. Go back to the start of the string and back to step 1
5. If the first symbol is a "+", remove the "+" and complete



Loop 1: 00+000
 Loop 2: 0+0000
 Loop 3: +00000 ^Q1
 Replace + with 0 ^Q5

Output: 00000

TM: Formally

- A transition $\delta(q_i, x) = (q_j, y, d)$ depends on:
 1. the current state q_i , and
 2. the current symbol x under the tape head
- A transition consists of three actions:
 1. change state to q_j
 2. over-write tape symbol x by y , and
 3. move the tape in direction $d = \{L, R\}$ (Left or Right)
 - If the tape head is on the leftmost symbol, moving left has no effect
- TM halts if it enters a state where there is **no move** – i.e., $\delta(q, x)$ is undefined

Outline

- Review
- Bell-LaPadula Multics interpretation
- General undecidability of security
 - Preliminaries
 - Halting problem
 - Turing machines
 - Safety question
 - Harrison-Ruzzo-Ullman result

What are we trying to do?

- Can we determine if a computer system is “secure” (i.e., implements our policy with high assurance)?
- Our goal is to answer the key question: is this protection problem decidable?



Objectives of the HRU Work


- Provide a model that is sufficiently powerful to encode several access control approaches, and precise enough so that security properties can be analyzed
- Introduce the “safety problem”
 - Accurately and concisely expresses the essence of the protection problem
- Show that the safety problem
 - is undecidable in general
 - is undecidable in monotonic case
 - A monotonic system is one in which rights cannot be deleted, and subjects and objects cannot be destroyed
 - is decidable in certain cases

Overview of the HRU Model

- What is “Secure”?
 - Adding a generic right r where there was not one is “leaking”
 - If a system S , beginning in initial state s_0 , cannot leak right r , it is **safe** with respect to the right r
- Safety Question
 - Does there exist an algorithm for determining whether a protection system S with initial state s_0 is safe with respect to a generic right r ?
 - Here, “safe” = “secure” for an abstract model
- Answer for a general case is “no”
 - Reduce halting problem to safety problem
 - Show that a Turing machine can be “modelled” by a protection system with the “states” of the machine mapped to the “rights” of the protection system

Notion of a “Leaked” Right

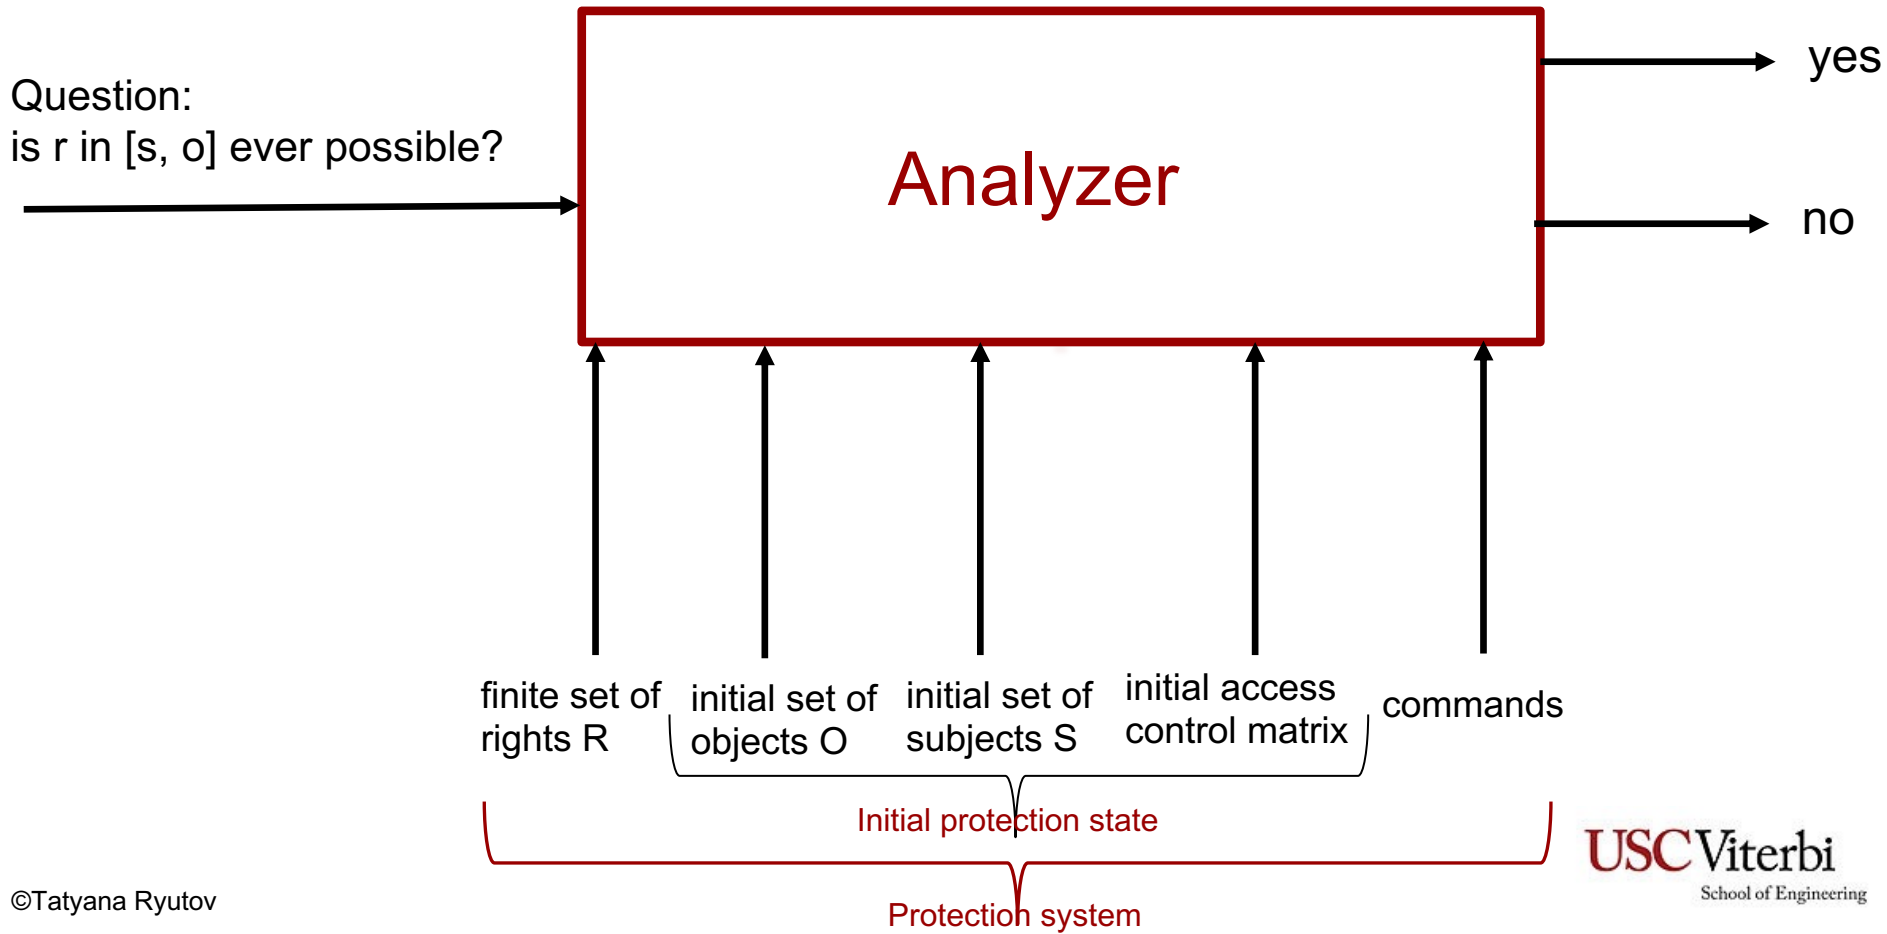
- Rights are the entries in access control matrix
 - Each subject has (or does not have) rights to an object
- Define **leaked**
 - Generic right added to element of access matrix when elements do not already contain the right
- Policy defines the authorized set of states
 - No command can leak a right r
- Define **safe** state with respect to a right r
 - System can never leak the right r
 - System is *unsafe* if it can enter unauthorized state



S_1 and S_2 - safe states, S_3 - unsafe state

HRU: General Idea

- Can we build an analyzer to answer the safety question?
 - Assume that the sole purpose of a system is to change access privileges to objects, ignoring other computations that might be occurring



HRU Protection System

- A protection system is a state-transition system
- A model for **protection** of computer system consists of:
 1. A finite set of generic rights R
 2. Initial protection state (initial set of objects O , initial set of subjects S , and initial access matrix P)
 3. A finite set of **commands** C of the form:

command $\alpha(X_1, X_2, \dots, X_k)$

if r_1 in (Xs_1, Xo_1) **and** r_2 in (Xs_2, Xo_2) **and** r_i in (Xs_i, Xo_i)

...

then

$op_1; op_2; \dots op_n$

end

parameters

conditions (test presence of certain rights in certain positions in access matrix)

command body is straight line code, no conditionals, no function invocation

- These commands are interpreted as a sequence of primitive **operations**

HRU Primitive Operations

- Primitive operations affect the state of access matrix:
 1. enter r into (X_s, X_o)
 - Condition: $X_s \in S$ and $X_o \in O$
 - r may already exist in (X_s, X_o)
 2. delete r from (X_s, X_o)
 - Condition: $X_s \in S$ and $X_o \in O$
 - r does not need to exist in (X_s, X_o)
 3. create subject X_s
 - Condition: $X_s \notin O$
 4. create object X_o
 - Condition: $X_o \notin O$
 5. destroy subject X_s
 - Condition: $X_s \in S$
 6. destroy object X_o
 - Condition: $X_o \in O$ and $X_o \notin S$

The State of A Protection System

- HRU define the “configuration” of a system
- Instantaneous description of protection system is a triple (S, O, P):
 - S is the set of current subjects,
 - O is the set of current objects, $S \subseteq O$ (subjects can be objects)
 - P is an access control matrix
 - one row for each subject
 - one column for each object
 - each cell contains a set of rights from R

How does state transition work?

- Given a protection system (R, C) , state q_1 can reach state q_2 IFF there is an instance of a command in C so that all conditions are true at state q_1 and executing the primitive operations one by one results in state q_2
 - R is a finite set of generic rights
 - C is a finite set of commands
- A command is executed as a whole (similar to a transaction), if one step fails, then nothing changes

} these sets do not change

Example 1

[Unix] process p creates file f with owner $read$ and $write$ (r, w) will be represented by the following:

Command $create_file(p, f)$

Create object f  create a column in the access matrix

Enter own into $a[p, f]$

Enter r into $a[p, f]$

Enter w into $a[p, f]$

End

Example 2

- Process p creates a new process q

Command *spawn_process*(p, q)

Create subject q ;  create a row in the access matrix

Enter *own* into $a[p,q]$

Enter r into $a[p,q]$

Enter w into $a[p,q]$

Enter r into $a[q,p]$

Enter w into $a[q,p]$

 parent and child can
signal each other

End

HRU: The Safety Problem

- Given a protection system and generic right r , we say that the initial configuration Q_0 is **unsafe** for r (or leaks r) if there is a configuration Q and a command α such that:
 - Q is reachable from Q_0
 - α leaks r from Q
- We say that a command $\alpha(x_1, \dots, x_k)$ **leaks** generic right r from Q if α , when run on Q , can execute a primitive operation which enters r into a cell of the access matrix which did not previously contain r

Relationship between TM and HRU

- It is undecidable (no generic algorithm) to determine whether an arbitrary TM halts or not
 - Or enters any arbitrary state q_f
- Idea: reduce protection problem (HRU) to TM
 - If TM enters state q_f , then the protection system can leak generic right r , otherwise, it is safe for r
 - Generic right r is arbitrary and hence yielding state q_f is also arbitrary
 - Since it is undecidable whether the TM enters arbitrary state q_f , it must be undecidable whether the protection system is safe for r
- Next question: how to map HRU to TM?

Mapping a Tape to an Access Matrix

- Create the protection matrix from the TM's tape
 - Encode the contents of the TM's tape on the diagonal of the protection matrix: element $[s, s]$ will contain the s^{th} tape square
 - How can we represent **sequential** tape?
 - Subject s_i represents cell i^{th} cell on the tape
 - Each subject s_i “owns” subject s_{i+1}
 - Sequential ownership relation represents **sequential tape**
 - Any cell (in TM) holding a symbol indicates subject s_i gave that right to itself
 - Last subject s_k has right *end*, indicating that subject s_{k+1} (which s_k owns), has not yet been created
 - The head is at the i^{th} cell and the current state is $q \Rightarrow q \in (s_i, s_i)$

	s1	s2	s3	...						
s1	A	own								
s2		B	own							
s3			C	own						
.				D	own					
.					E	own				
.						F	own			
.										
.										
.										
.										

A	B	C	D	E	F	G	H	
---	---	---	---	---	---	---	---	--

Mapping a Tape to an Access Matrix


- We also need to encode the state of the TM
 - The set of generic rights represent states and tape symbols
 - Two special rights: *own* and *end*
 - *end* is the last cell before blanks
 - For example, rights *g* and *c* are elements of $[s_3, s_3]$ when the TM is in state *q*, the read/write head is on square 3, and symbol *c* is in $[s_3, s_3]$
- Turing Machine instructions are mapped to commands of the HRU protection system

Representing Head Moves

- The moves of TM are represented as HRU commands
 - Changing the state of the TM is equivalent to commands that delete and add rights, objects, and subjects
- Choose HRU commands to represent TM moves
 - Right: $\delta(k, C) = (k_1, X, R)$
 - E.g., if cell k is the current position, command $(k_1, C/X, R)$ substitutes access right C for access right X in the cell, and k_1 is the cell to the immediate right of k
 - When move right till end (blanks), need to create a new subject
 - Left: $\delta(k, C) = (k_1, X, L)$
 - E.g., if cell k is the current position, command $(k_1, C/X, L)$ substitutes access right C for access right X in the cell, and k_1 is the cell to the immediate left of k
- For any possible TM transition, can have corresponding HRU command

Mapping Depiction

TM



Current state is k

Current symbol is C

head

HRU Matrix

	s_1	s_2	s_3	s_4	
s_1	A	own			
s_2		B	own		
s_3			$C\ k$	own	
s_4				D end	

Symbols, States \Rightarrow rights

Tape cell \Rightarrow subject

Cell s_i has A $\Rightarrow s_i$ has A right on itself

Cell $s_k \Rightarrow s_k$ has end right on itself

State p , head at $s_i \Rightarrow s_i$ has p right on itself

Distinguished right own: s_i owns s_{i+1} for $1 \leq i < k$

Right Move (Left is Symmetrical)

```

command  $c_{k,c}(s_3, s_4)$ 
if own in  $A[s_3, s_4]$  and  $k$  in  $A[s_3, s_3]$ 
    and  $C$  in  $A[s_3, s_3]$ 
then
    delete  $k$  from  $A[s_3, s_3]$ ;
    delete  $C$  from  $A[s_3, s_3]$ ;
    enter  $X$  into  $A[s_3, s_3]$ ;
    enter  $k_1$  into  $A[s_4, s_4]$ ;
end
    
```

1	2	3	4	
A	B	C	D	

Current state is k

Current symbol is C




head

	s_1	s_2	s_3	s_4	
s_1	A	<i>own</i>			
s_2		B	<i>own</i>		
s_3			C k	<i>own</i>	
s_4				D <i>end</i>	

$$\delta(k, C) = (k_1, X, R)$$

After One Right Move



Current state is k_1

Current symbol is C

↑
head

$$\delta(k, C) = (k_1, X, R)$$

	s_1	s_2	s_3	s_4	
s_1	A	own			
s_2		B	own		
s_3			X	own	
s_4				D k_1 end	

Right Move at End

```

command crightmostk,c(s4, s5)
if end in A[s4, s4] and k1 in A[s4, s4]
    and D in A[s4, s4]
then
    delete end from A[s4, s4];
    create subject s5;
    enter own into A[s4, s5];
    enter end into A[s5, s5];
    delete k1 from A[s4, s4];
    delete D from A[s4, s4];
    enter Y into A[s4, s4];
    enter k2 into A[s5, s5];
end

```

1	2	3	4	
A	B	X	D	...

Current state is k_1


Current symbol is D

↑
head

$$\delta(k_1, D) = (k_2, Y, R)$$

	s ₁	s ₂	s ₃	s ₄	
s ₁	A	own			
s ₂		B	own		
s ₃			X	own	
s ₄				D k_1 end	

After Right Move at End



Current state is k_1

Current symbol is D

↑
head

	s_1	s_2	s_3	s_4	s_5
s_1	A	own			
s_2		B	own		
s_3			X	own	
s_4				Y	own
s_5					k_2 end

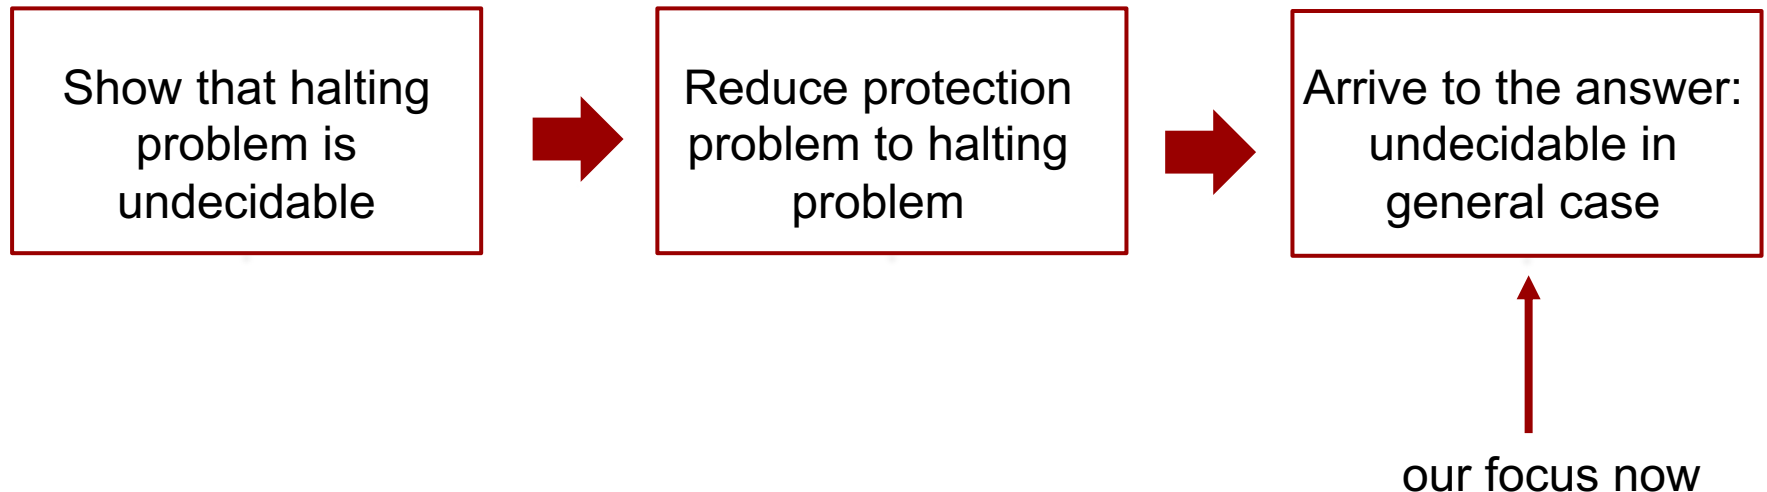
$$\delta(k_1, D) = (k_2, Y, R)$$

Rest of Proof

- Protection system exactly simulates actions of TM
 - Exactly one *end* right – blanks after that
 - Exactly one access right corresponds to a TM state
 - Thus, at most one applicable command
- If TM enters state q_f , then right is leaked
- Generic right r is **arbitrary** and hence yielding state q_f is also arbitrary
- If safety question is decidable, then we can represent TM as discussed and determine if q_f leaks
 - Implies halting problem is decidable. **A contradiction!**
- Conclusion: safety problem **undecidable**

What are we trying to do?

- Can we determine if a computer system is “secure” (i.e., implements our policy with high assurance)?
- Our goal is to answer the key question: is this protection problem decidable?



Theoretical Limits on System Security Summary

- Harrison, Ruzzo and Ullman (HRU) defined "safety" problem for protection systems
 - Safety refers to some **abstract** model
 - Security refers to **actual** implementation
- System can only be secure if it implements a policy based on a **safe model**
 - But a safe model does NOT ensure a secure system
- DAC has fundamental flow control limitation
 - It is generally unsafe model
 - Consistent with analysis of Trojan horse threat
- MAC can be safe
 - Imposes sufficient restrictions on right propagation
- How can we use this in practice?
 - Balanced assurance

